

Seeding metaheuristics with learned and heuristic solution pools for the Interval Job Shop Scheduling Problem: an anytime study of when initialisation helps

Hernán Díaz Rodríguez^a

^a*Department of Computer Science, University of Oviedo, Spain*

Abstract

Practitioners routinely assume that seeding a metaheuristic’s initial population with good solutions improves its results. We test that on the Job Shop Scheduling Problem with interval processing times, using four published solvers spanning the weak-to-strong axis, four seed generators — a deep reinforcement-learning policy and a ladder of randomised constructives — and mixed pools, over 61 interval Taillard instances and 30 runs per cell (58,560 runs). A pilot study first identifies three practices that manufacture positive seeding results: stagnation-triggered stopping, single-budget comparison, and self-referential speed metrics.

The answer turns almost entirely on the budget. At the full budget, which a plateau rule finds sufficient for convergence on only 133 of the 244 control cells, seeding is worth 0.1% to 1.0% of makespan and is significant on one solver of four; at a tenth of it, 1.3% to 5.2% on all four. Seeding buys time rather than quality: it reaches a given quality sooner, not a better one — except on the weakest solver, where it does both.

Two design levers are separable. Composition rather than seed quality orders the pools that help: a mixed pool starts from a worse average population than its components and finishes ahead of them. Independently, the fraction seeded interacts with the solver, half the population being better on the genetic algorithm and worse on the elitist ABC. An exploratory eighth arm combining

Email address: `herdiaz21@uniovi.es` (Hernán Díaz Rodríguez)

both is the best performer on three solvers and the only one that improves the genetic algorithm. Finally, the premature convergence usually blamed for seeding failures is contradicted: the harmed genetic algorithm ends *more* dispersed than its control.

Keywords: Job shop scheduling, interval uncertainty, population seeding, metaheuristics, reinforcement learning, anytime behaviour

1. Introduction

Real-world scheduling rarely enjoys exact knowledge of task durations. When only lower and upper bounds are available, the Job Shop Scheduling Problem with interval processing times (IJSP) models each duration as a closed interval and propagates interval arithmetic through the schedule [1, 2]. The IJSP has accumulated a substantial algorithmic literature over the last years, including genetic algorithms [3], elitist artificial bee colony variants [4, 5], fast memetic ABC hybrids [6], and, most recently, tabu-search hybrids built on extreme-critical-path neighbourhoods that constitute the current state of the art [7].

In parallel, learning-based methods have begun to produce fast constructive policies for shop scheduling [8, 9]. A deep reinforcement-learning (RL) dispatching policy for the IJSP can emit thousands of feasible, above-average solutions per instance in seconds, raising an obvious question: should we seed the initial populations of existing metaheuristics with such solutions? Folk wisdom answers yes, and heuristic seeding of evolutionary algorithms has a long history [10, 11]. Yet the empirical evidence behind that wisdom is weaker than it appears. Seeding studies typically compare final objective values under a fixed stopping rule, and a pilot study reported in Section 5 shows that practice to be fragile in three specific ways — *stagnation-triggered stopping*, *budget-dependent gains that reverse*, and *self-referential speed metrics* — any one of which can manufacture a positive seeding result where none exists. The design of this study follows from avoiding them.

This paper makes three contributions.

1. **A pre-specified anytime methodology** for evaluating population seeding: convergence-checked final quality; two complementary time-to-target metrics (time to the arm’s own plateau, and time to the unseeded control’s final quality) reported with the proportion of instances on which the target is reached; and paired runs with per-cell budgets calibrated on the control.
2. **A portfolio study across the weak-to-strong axis** — GA [3], elitist ABC [4], memetic fEABC-LS [6], and the tabu hybrid TS-N2 [7] — showing that the sign of the seeding effect is *not* predicted by solver strength: the weakest solver gains, the strongest gains on the largest instances, one is unaffected, and one is degraded. The sign tracks whether the solver’s variation operators can consolidate a structurally heterogeneous population, and the magnitude tracks the ratio of instance size to search budget; both are associations our design supports rather than causal effects it isolates.
3. **Learned versus heuristic seeds:** the RL policy is compared, under identical injection protocols, against a quality ladder of randomised constructives (machine-operations-remaining, randomised Giffler–Thompson, and a tuned genetic-programming rule), plus a mixed pool assembled from the same solutions. The mixed pool starts from a *worse* average population than its single-generator counterparts and still converges better, which points to pool composition rather than seed quality as the operative variable.

Our results do not support the assumption we set out to test. A mixed seed pool improves the weakest solver decisively, leaves a memetic ABC unchanged, does not harm a genetic algorithm although every single-generator pool does, and helps a strong tabu hybrid only where its budget is tightest relative to instance size. The sign of the effect is therefore not predicted by solver quality; what it tracks is whether the solver’s variation operators can exploit a structurally heterogeneous population, and its magnitude tracks the ratio of instance size to search budget. Section 7 gives the figures.

The paper is organised as follows. Section 2 reviews related work. Section 3 defines the IJSP and the solver portfolio. Section 4 describes the seed generators, the pool audit and the injection mechanism. Section 5 reports the pilot study whose failure modes motivate the design of Section 6. Section 7 presents the results, Section 8 discusses their implications, and Section 9 concludes.

2. Related work

From the deterministic job shop to the interval job shop.. The solvers studied here inherit their machinery from the deterministic job shop. Simulated annealing over the disjunctive graph [12] and the branch-and-bound of [13] established the representation; the decisive step for local search was the critical-block neighbourhood of the fast taboo search of Nowicki and Smutnicki [14], later refined in [15], and the guided local search of [16]. The block-boundary neighbourhood used by the tabu component of our strongest solver is the interval counterpart of that construction, and the local search of [17] carried it to uncertain durations.

Metaheuristics for the IJSP.. Lei introduced population-based neighbourhood search for interval makespan minimisation [1] and a GA for interval due-date objectives [2]. A GA tailored to interval uncertainty with an insertion schedule-generation scheme was proposed in [3], later outperformed by elitist ABC variants [4], a seasonal elitist ABC [5], and the fast elitist ABC with hill-climbing local search (fEABC-LS) [6]. The most recent entry couples the fEABC framework with a best-improvement tabu search over block-boundary neighbourhoods defined on both extreme scenario graphs [7]. Our portfolio comprises exactly these published solvers, replicated from their original configurations (Section 3). The neighbouring fuzzy formulation has received sustained attention to benchmarking in particular: test beds, models and lower bounds for the fuzzy job shop are given in [18, 19], and we follow that practice in reporting deviations against published bounds rather than against best-known values of our own (Section 7.1).

Population initialisation and seeding. Initialisation is one of the oldest levers in evolutionary computation and one of the least settled. Surveys of the mechanisms are available in [10, 20, 21]; the recurring families are quasi-random and space-filling designs [22], opposition-based sampling [23], constructive-heuristic seeding, and the injection of solutions retained from previous runs, studied as case-injected genetic search [11]. Two findings from this literature frame the present study. First, seeding is not uniformly beneficial: Burke et al. observed in evolutionary timetabling that when the seeding procedure returns similar solutions, the resulting loss of diversity can leave the algorithm worse off than random initialisation [24], and computational studies in the multi-objective setting report that the benefit varies sharply across problem families [25] and that injecting even optimal solutions is not automatically an improvement [26, 27]. Second, the evidence base is thin where it matters: the review of [21] concludes that published initialisation studies rarely establish *why* a scheme works or *under which conditions*, and that the interaction between initialisation and the specific algorithm being initialised remains under-examined. What distinguishes the present study is therefore not the mechanism — ours is deliberately the simplest possible: replace k random individuals by pool solutions at initialisation — but that the same injection protocol is applied across a portfolio of solvers of differing strength under a common budget, with convergence checks, anytime metrics and paired statistics with multiple-comparison control.

Learning to construct schedules. Deep RL dispatchers for the deterministic JSP [8, 28], their extension to the flexible variant [29], and the broader learn-to-optimize programme [9, 30] produce fast constructive policies. Genetic programming pursues the same end by evolving dispatching rules rather than learning them by gradient, and the two lines are surveyed together in [31]. Where these policies meet mature metaheuristics, the coupling is usually made at the level of control — for example, using a learned agent to adapt operators or parameters during the search [32]. The alternative coupling is initialisation, which leaves the solver untouched and is therefore the one that isolates the value of the

learned solutions themselves. We treat the learned policy as one seed generator among several, which allows the question *is a learned constructive better than a tuned heuristic one as a seed source?* to be answered under identical injection protocols.

Experimental methodology... Because a seeded run and its control differ from the first generation, comparing them only at the end of the budget discards most of the effect. We follow the anytime perspective of [33], reporting time-to-target alongside final quality, and use non-parametric paired procedures with multiple-comparison control [34] and a stochastic effect size [35]. No parameter of any solver is tuned in this study: configurations are taken verbatim from the source publications, so that automatic configuration [36] is deliberately excluded as a source of variation between arms.

3. The IJSP and the solver portfolio

3.1. Problem definition

An IJSP instance consists of n jobs, each a fixed sequence of m operations, one per machine. The processing time of operation o is a closed interval $p_o = [p_o^-, p_o^+]$. Interval addition and component-wise maximum extend the classical recursion for starting and completion times: for a feasible processing order σ , the starting time of o is $s_o = \max(c_{PJ(o)}, c_{PM(o)})$ with $c_o = s_o + p_o$, where PJ and PM denote job and machine predecessors and the maximum is taken component-wise [1, 3]. The makespan $C_{\max} = [C_{\max}^-, C_{\max}^+]$ is the component-wise maximum of job completion times; it is exact, because the makespan is monotone in the durations, so the all-lower and all-upper configurations realise its bounds. Comparing schedules requires an interval ranking operator; we use the expected-value ranking of [37], i.e. the interval midpoint, which we write $\mathbb{E}[C_{\max}]$ following the convention of this literature; no probability distribution over the interval is assumed, and the quantity is a ranking criterion rather than a mathematical expectation. It ranks solutions for the evolutionary solvers, following their original papers, while LEX2 (lexicographic-by-upper-bound) is

used inside TS-N2, following [7]; all cross-arm comparisons are made on $\mathbb{E}[C_{\max}]$ within a single solver, so the ranking choice never confounds the seeding effect.

Instances are the standard interval Taillard benchmark: each crisp duration p_o becomes $[p_o - \delta, p_o + \delta]$ with δ drawn uniformly from $[0, 0.15 p_o]$ [3, 38].

3.2. The solver portfolio

All four solvers share the permutation-with-repetition encoding [39], the insertion schedule-generation scheme (active schedules) for decoding, population size 250, and the component-wise interval evaluation described above. Table 1 summarises their published configurations, whose parameter settings we replicate verbatim.

3.3. Replication check

Before any seeding experiment we replicated each solver against its published results on the twelve classical interval instances. The GA and fEABC replications match the published average relative errors within ~ 1 percentage point on all twelve instances (fEABC matches or slightly improves on every one); fEABC-LS reproduces the published ranking (best of its family). Our ABCE3 replication runs systematically ~ 1.5 – 2.5 pp above the published averages; we were unable to close this gap from the published description and disclose it here.

We take the published parameter settings verbatim, but that is a statement about the configuration, not a claim that our implementation is behaviourally identical to the authors’ own. The ABCE3 gap shows it need not be, and the local-search footnote to Table 1 shows where the two can diverge while the parameter file reads the same.

No conclusion of this paper rests on matching the published figures, because every comparison is within-solver: a seeded arm against the unseeded control of the *same* implementation, under the same budget, on the same instances. Any shortfall is shared by both sides of every contrast.

One consequence does deserve stating, because it bears on how far the results generalise. ABCE3 is at once the weakest control in the portfolio, the solver

Table 1: Solver portfolio (published configurations). All use population 250, permutation-with-repetition encoding and insertion SGS decoding.

	GA [3]	ABCE3 [4]	FEABC-LS [6]	TS-N2 [7]
Frame	generational GA	elitist ABC	fast elitist ABC	fast elitist ABC
Crossover	JOX, $p_c=1.0$	GOX, $p=1.0$	JOX	JOX
Mutation/move	insertion, $p_m=0.15$	swap, $p=1.0$	swap	insertion
Elite / trials	—	$B=50$, $NT_{\max}=20$	40 / 15	50 / 24
Local search	none	none	HC (first impr.), N_1	tabu (best impr.), N_2 , LEX2
LS control	—	—	best of each trial pair [†]	100%, $B_{\max}=20$, $T_{\max}=2s$

[†] The published configuration sets the local-search target to 75% of the population, and that is the value we ran. Its effect in this implementation is narrower than the figure suggests, and we state what the code does rather than what the parameter names: local search is invoked on the two-member candidate population formed by each employed-bee trial, and $\lfloor 0.75 \times 2 \rfloor - 1 = 0$ leaves the random quota empty, so only the better of the two candidates is improved (`ArtificialBeeColonyPS0.cpp:559–574, 705, 727`). Every arm of this solver runs the identical mechanism, so within-solver comparisons are unaffected; the qualification bears on how the solver should be described, and on the replication gap reported below.

seeding helps most, and the one whose replication falls short. The shortfall, 1.5–2.5 points, is larger than the 1.14-point gain seeding produces there. The within-solver contrast remains valid, but the headroom that seeding exploits on this solver may be headroom our implementation left, not headroom intrinsic to the algorithm. Read the finding as “seeding helps a solver that leaves this much on the table”, which is what we measured, rather than as a claim about published ABCE3.

4. Seed pools and injection mechanism

4.1. Generators and pools

For every instance we use pools of 1,024 i.i.d. solutions per generator, each stored as a job-permutation plus its makespan interval under a semi-active decoder. Four generators span a quality ladder:

- **MOR- ε** : most-operations-remaining dispatching with $\varepsilon = 0.1$ random tie-perturbation (weak baseline);
- **GT- ε** : randomised Giffler–Thompson with ε -noise inside the conflict set (active schedules; strong baseline);
- **GP- ε** : a genetic-programming priority rule, evolved and irace-tuned offline, with GRASP-style ε -noise (strongest heuristic baseline);
- **v2**: a deep-RL dispatching policy (deep-sets architecture, PPO, stochastic sampling of the policy for diversity), trained only on Taillard instances TA11–TA14 and developed on TA15–TA20. Its pool is the one structural exception to i.i.d. sampling: it is three consecutive blocks of 342/341/341 solutions, one per training checkpoint, and the blocks differ measurably in quality (Kruskal–Wallis $p < 10^{-6}$ on every instance tested). The consequence for the arms that use it is quantified in Section 8.4. The policy also ranks candidate schedules by LEX2, whereas three of the four solvers rank by the midpoint, so these seeds are optimised under a different criterion from the one that later evaluates them.

Across the 71-instance pool collection the ladder is clean: mean relative errors of 51.1% (MOR) > 30.3% (GT) > 21.4% (GP) $\approx$ 21.2% (v2), with v2 exhibiting the highest structural diversity of the four (0.95 vs. 0.92 for GP on a normalised disjunctive-order distance) and GP slightly stronger on the largest size classes. A fifth, mixed pool (equal parts v2/GT/GP; MOR is deliberately excluded, being the weakest generator) is used by the *mix* arm. All five pools — the four

generators and the mixture — contain 1,024 solutions per instance, so no arm is sampled more richly than another.

The mixed pool interleaves the three generators one entry at a time, so that any contiguous window of 250 entries contains 84/83/83 seeds of v2/GT/GP to within one seed, whatever its offset.

4.2. Evaluation convention of the pools

Each pool line stores a permutation together with its interval makespan, evaluated component-wise so that stored and solver-side objective values coincide. The distinction matters because both interval conventions of Section 3 occur in this literature and they do not agree: against LEX2 aggregation the upper bound is identical on all 1,024 lines of every pool, while the lower bound differs on 1–22% of lines, by up to 98 time units on the 50×20 instances. Anyone exchanging solutions between components that rank intervals differently should expect the same.

4.3. Injection mechanism

Seeding is deliberately minimal. A seeded arm with fraction $k/250$ replaces the first k individuals of the initial population by pool lines; the remaining $250 - k$ individuals are generated by the solver’s standard random initialisation. For paired run $r \in \{0, \dots, 29\}$ the arm takes pool lines $[(r \cdot k) \bmod 1024, (r \cdot k + k) \bmod 1024)$ — consecutive blocks, no reordering or filtering, so each run draws from the generator’s true output distribution rather than from a curated subset. The 30 blocks start at 30 distinct offsets and consecutive ones are disjoint, but with $k = 250$ and 1,024 pool lines the wraparound makes any two of them share 22% of their entries on average, and each pool line is drawn by 7.3 of the 30 runs; the seeded populations are therefore not fully independent across runs, which is one reason the unit of analysis throughout is the instance rather than the run. Scout-bee replacements during the run remain fully random. Seeds enter as raw permutations and are re-decoded and re-evaluated by the receiving solver, so no foreign objective values ever enter the search.

4.4. The cost of producing the seeds

Seeding is not free, and at the budgets where it helps most the cost is not negligible. Generation was timed for the two heuristic generators on a single thread: a pool of 1,024 solutions takes 91 s on 15×15 and 1,255 s on 50×20 for GP- ε , and 89 s to 1,500 s for GT- ε , scaling roughly linearly in the number of operations. The learned policy is the expensive one: measured per-rollout inference of 0.60 s on 15×15 and 3.98 s on 50×20 puts it six to ten times above the rules per solution. We did not time the v2 and MOR- ε pools directly and do not report an equivalent figure for them.

Two readings follow, and they differ by an order of magnitude. Producing the 250 seeds of a *single* run costs from 22 s on 15×15 to 306 s on 50×20 with the cheapest generator — between four and fifty-one times the tenth of the budget at which the anytime advantage is largest, depending on solver and class. Amortised instead over the thirty runs that share a pool, as in this study and in any repeated-solve setting, the same cost falls to between 0.1 and 7 times that slice, and below it in most cells. The short-budget result therefore holds as stated for a practitioner who reuses a pool, and does not hold for one who generates seeds to solve a single instance once.

5. Pilot study: three ways to fool yourself

We first ran a full pilot on three clean instances (tai15_15_05, tai20_20_02, tai30_20_04; lower bounds 1224, 1581 and 1952) with six solvers (the portfolio plus two isolation variants), seven arms and 30 paired runs per cell. The pilot’s role in this paper is methodological: each of its three phases exposed an artefact that the final design of Section 6 removes.

5.1. Artefact 1: stagnation-triggered stopping

Under the customary rule of 25 consecutive non-improving generations — the stopping criterion of every published solver in the portfolio — seeded GA arms terminated at the 25-generation floor (the solver could not improve the injected

solutions), while the unseeded control ran for 100–150 generations. The resulting “final” comparison charges the seeds with the solver’s inability to improve them. Any seeding study that uses stagnation-based stopping inherits this bias.

5.2. Artefact 2: budget-dependent gains that reverse

Re-running everything under a fixed budget of 200 generations (no stagnation rule) appeared to settle matters: seeded GA arms beat the control on the hardest instance by 34 units of expected makespan. A convergence check, however, showed the GA control still improving at budget exhaustion. At 5× the budget (1,000 generations) the control kept descending, overtook the stagnant seeded populations, and the effect reversed sign (from −34 to +67 on tai30_20_04 with full v2 seeding). The same check showed the elitist ABC’s seeded advantage *persisting* at convergence (−46 to −35 with GP seeds), and the strong local-search solvers saturating at equal quality. A separate isolation experiment — attaching the identical tabu local search of TS-N2 to a GA frame — initially suggested the GA frame benefited from seeds where the ABC frame did not; extending its time cap dissolved that difference too, tracing it to budget starvation rather than to the frame. *Conclusion: no seeding claim about final quality is meaningful without a convergence check of the unseeded control.*

5.3. Artefact 3: what “faster” means

We measure convergence speed with two deliberately different metrics: (A) time to reach within 1% of the *arm’s own* final value, and (B) time to reach the *unseeded control’s* final quality (tolerance +0.5%). Metric A is self-referential: an arm that settles quickly on a worse plateau scores well, so an apparent acceleration may be nothing more than earlier stagnation. Metric B removes that degree of freedom by fixing a common target, at the cost of being undefined on instances where the arm never reaches it — which makes the proportion of instances reaching the target as much a part of the result as the time itself. Reporting both, with that proportion, prevents the most common over-claim in anytime comparisons.

5.4. *What survives the artefacts*

With all three corrections in place, the pilot’s surviving picture is: (i) seeding improves converged quality for the elitist ABC on all three instances (Wilcoxon, paired); (ii) for solvers with strong local search the final quality saturates but metric-A acceleration remains; (iii) for the plain GA the apparent gain is an artefact; (iv) GP and v2 seeds are statistically tied at the top of the generator ladder, MOR seeds never help; and (v) initial-population quality ranks exactly as the pool ladder predicts. These five statements were fixed before the final experiment and are what it tests at scale; Section 7.11 reports which of them survived.

6. Final experimental design

6.1. *Instances*

All 60 interval Taillard instances TA1–TA70 excluding TA11–TA20, plus ft10 (61 total). TA11–TA14 are the RL training set and TA15–TA20 the RL/GP development set; excluding them is mandatory because they contaminate the *seed generators*. Twenty instances used for irace tuning of TS-N2’s hyper-parameters are retained but flagged: solver-tuning contamination applies equally to every arm of that solver and cancels in within-solver paired comparisons; a sensitivity analysis replicates all conclusions with the flagged instances removed. Size classes range from 15×15 to 50×20 (225–1,000 operations). Taillard’s largest class, TA71–TA80 (100×20 , 2,000 operations), is also excluded: at twice the operation count of our largest class it falls outside the budget calibration of Section 6, and extending the design to it would require a separate calibration that we did not perform. Consequently no seed pools were generated for either excluded group.

6.2. *Arms*

Seven arms per solver were fixed before the experiment ran: the unseeded control A0; the generator ladder MOR, GT, GP, v2 and mix at the pre-specified

fraction p^* ; and a v2@50% contrast arm. An eighth, mix@50%, was added afterwards as an exploratory follow-up and is reported separately in Section 7.8; it ran under the identical protocol, bringing the design to $61 \times 8 \times 4 \times 30$ cells. The fraction sub-study (Phase 0, pre-specified decision rule combining converged quality in ABCE3 and time-to-target in TS-N2 over eight stratified instances) selected $p^* = 100\%$: the benefit of v2 seeding is monotone in the fraction on both criteria, with no interior optimum.

6.3. Budgets and stopping

Budgets are set per (*solver, size class*) rather than per instance. For each class we take one or two calibration instances, run the unseeded control on them, and record t_{conv} : the first CPU time at which the run-averaged trace reaches within 0.1% of its own final value. The budget applied to every cell of that class is

$$B(\text{solver}, \text{class}) = \max\{60 \text{ s}, \min\{1.5 \tilde{t}_{\text{conv}}, 900 \text{ s}\}\}, \quad (1)$$

where $\tilde{t}_{\text{conv}}$ is the median over that class’s calibration instances. The floor of 60 s prevents degenerate budgets on the small classes, where several solvers plateau in a few seconds. Every arm of a cell receives the identical budget, which is what the within-solver comparisons require. The eight calibration instances — tai15_15_01, tai20_20_02, tai30_15_01, tai30_20_04, tai50_15_01, tai50_15_05, tai50_20_01 and tai50_20_05 — are stratified across the size classes and are themselves members of the 61-instance evaluation set; the control traces behind t_{conv} comprise five runs for the GA and fEABC-LS and thirty for ABCE3 and TS-N2. The horizon of a class is therefore the control’s own convergence timescale, measured on one or two of the instances to which it is later applied. What the endpoint measures is thus performance at a control-calibrated horizon rather than a solver-neutral approximation to convergence, a distinction we return to in Section 8.4.

The cap of 900 s is a feasibility constraint rather than a methodological choice. The design comprises 58,560 runs; at the budgets of (1) it consumed

roughly 4,900 CPU-hours and fourteen days of wall-clock time on fourteen concurrent workers. Raising the cap would scale that figure proportionally on the classes that reach it, which are also the slowest. We therefore accept the cap and treat its consequence explicitly: on a substantial fraction of cells the unseeded control is still improving when the budget expires, so those cells cannot support a converged-quality claim. Section 7.2 identifies them and reports the final-quality results both over all instances and restricted to the cells that do satisfy the plateau rule.

6.4. Statistics

30 runs per cell. Run r of every arm of a solver is launched with the same RNG seed index, which blocks the design by (instance, run index). This is *not* a common-random-numbers coupling: a seeded arm does not draw from the generator to build its first k individuals, so the generator state diverges between arms immediately after initialisation and the runs follow unrelated random trajectories thereafter. The pairing is therefore a labelling of blocks rather than a variance-reduction device. Signed-rank tests on such pairs remain valid — under the null the per-instance differences are still symmetric about zero — but they are conservative rather than more powerful, and we make no claim of matched trajectories. Within-solver comparisons use paired Wilcoxon signed-rank tests with Holm correction per family, plus Vargha–Delaney $\hat{A}_{12}$ effect sizes [35, 34]; generator rankings use Friedman tests with post-hoc analysis. Anytime results are reported as expected-makespan-versus-CPU-time curves and as the two time-to-target metrics of Section 5, the latter accompanied by the proportion of instances on which the target is reached — which carries the information a success-rate curve would convey at the budget limit.

7. Results

All results below come from the full factorial design of Section 6: 61 instances $\times$ 8 arms $\times$ 4 solvers $\times$ 30 RNG-paired runs (58,560 runs). Every instance–arm cell contains exactly 30 runs; no cell was imputed or dropped.

7.1. Reference for relative deviations

Relative percentage deviation (RPD) is measured against published lower bounds rather than against best-known values produced by our own group, which would bias RPD downwards. Our interval instances are obtained from Taillard’s benchmark [38] by a symmetric perturbation of each duration; we verified across all 57,450 operations that $\text{mid}(p_{ij}) = c_{ij}$ exactly and that routings are unchanged. Because the makespan is composed only of additions and maxima, and because $\text{mid}(A+B) = \text{mid}(A) + \text{mid}(B)$ while $\text{mid}(\max(A, B)) \geq \max(\text{mid}(A), \text{mid}(B))$ under the componentwise maximum, induction over the schedule DAG yields $\text{mid}(C_{\max}^{\text{int}}(x)) \geq C_{\max}^{\text{crisp}}(x)$ for *every* schedule x , and therefore

$$\mathbb{E}[C_{\max}]^* \geq C_{\max}^{\text{crisp}*} \geq \text{LB}, \quad (2)$$

so the crisp lower bound is a valid (if not tight) bound for the interval problem. We take LB from the consolidated table of Coupvent des Graviers et al. [40]; the instance mapping was checked rather than assumed, with 19 of 60 instances showing exact agreement between the published bound and the machine-load bound recomputed from the instance files. For the remaining 41 the check is an inequality, which is necessary but does not identify an instance, and we do not claim more than that.

An exact identification is nevertheless available to any reader, and we supply what it needs. Because each interval was built as $[p_o - \delta, p_o + \delta]$, it is symmetric about the original crisp duration, so the midpoint recovers that duration exactly; across the 61 instances every midpoint is integral, with no exception. We therefore release the recovered crisp matrices, one per instance with a hash, so that identification reduces to an element-wise comparison against Taillard’s published files. The recovery procedure itself is validated on the one instance of the set whose crisp form is unambiguous: `ft10` recovers to the Fisher–Thompson 10×10 matrix exactly.

Inequality (2) also yields a falsifiable prediction: no run may ever fall below its bound. Across all 58,560 runs there is not a single violation. This is a

consistency check rather than a proof of correctness: many errors would preserve the inequality, and it is the released crisp matrices, not this check, that would settle the mapping. Moreover the bound is *attained exactly* on six instances — `tai30_15_05`, `tai50_15_01`, `tai50_15_02`, `tai50_15_04`, `tai50_15_07` and `tai50_20_09` — where our portfolio reaches $\mathbb{E}[C_{\max}]$ equal to the proven crisp optimum. By (2) these schedules are therefore optimal for the interval instance as well, which both certifies optimality on those six instances and confirms that the inequality is tight rather than vacuous.

7.2. Which cells converged

The plateau rule of Section 6 — improvement in $\mathbb{E}[C_{\max}]$ below 0.1% over the final tenth of the budget — is applied to each control *run* against the solution the solver returns, rather than to the run-averaged curve, since averaging thirty traces can hide a subset that is still improving; a cell counts as converged when at least 90% of its runs satisfy it. This classifies 111 of the 244 control cells as *anytime-only*: 49 of 61 for the GA, 25 for fEABCLS, 22 for TSN2 and 15 for ABCE3. The GA fails on every instance of the four largest classes, and the 50×20 class defeats every solver but ABCE3, whose own failures are otherwise confined to 15×15 , where a third of its runs are still improving when the 60 s floor stops them. What the classification is and is not sensitive to is reported in Section 7.12; the per-cell classifications are released with the data.

Table 2: Mean RPD (%) over the n instances whose *control* satisfies the plateau rule. Best arm per row in bold; compare Table 4 for the full 61.

Solver	n	A0	V2H	V2	MOR	GT	GP	MIX	MIXH
GA	12	4.75	4.28	4.39	6.82	5.35	4.49	4.00	4.13
ABCE3	46	11.61	10.46	10.34	11.85	11.28	10.72	10.35	10.43
fEABCLS	36	3.39	3.34	3.34	4.80	4.07	3.73	3.40	3.37
TSN2	39	2.75	2.82	2.83	3.54	3.41	3.07	2.78	2.79

Table 2 shows that the restriction leaves two of the four solvers unchanged in every respect that matters. fEABCLS keeps the same best arm and the

Table 3: Mean RPD (%) restricted to the instances where the control *and* the compared arm both satisfy the plateau rule. Negative Δ favours the seeded arm. Three representative arms; the full set is released with the data.

Solver	Arm	n both	A0	Arm	Δ
GA	V2	11	4.44	4.04	-0.40
	MOR	8	3.88	5.98	+2.10
	MIX	10	4.14	3.45	-0.69
ABCE3	V2	46	11.61	10.34	-1.27
	MOR	43	11.69	11.95	+0.26
	MIX	46	11.61	10.35	-1.26
fEABCLS	V2	31	3.26	3.22	-0.05
	MOR	19	3.86	5.27	+1.41
	MIX	34	3.46	3.48	+0.02
TSN2	V2	30	2.49	2.57	+0.08
	MOR	33	2.80	3.57	+0.78
	MIX	35	2.32	2.34	+0.02

same ordering. ABCE3 keeps its conclusion — V2 and MIX are separated by a hundredth of a point and both beat the control by more than a point — while its absolute level rises, because the cells it fails are the small 15×15 ones on which it does best.

Because the rule is evaluated on the control alone, the sound comparison is the one restricted to instances where *both* the control and the arm being compared plateau. Convergence rates differ enough between arms to make this matter: on fEABCLS the mixed pool plateaus on 44 of 61 instances against the control’s 36, while on the GA the randomised Giffler–Thompson pool plateaus on only 8. Table 3 reports those paired intersections.

The GA is the first exception, and the difference is not cosmetic. Over all 61 instances the unseeded control is the best arm (7.52 against MIX’s 7.58); over the 10 instances where control and mixed pool both plateau, the ordering

reverses and the mixed pool wins by 0.69 points (3.45 against 4.14). The two other pools that help at convergence, V2H and GP, reverse by a similar margin, while MOR and GT remain harmful there (+2.10 and +1.28). The negative result for the GA is therefore a statement about *insufficient budget*, not about seeding as such: it is the solver that needs the most time to consolidate a seeded population, and on the cells where it is given that time the harm disappears for the pools that were never the problem. This revises the reading of Section 8.2: what the GA lacks is not compatibility with heterogeneous parents in principle, but enough generations to exploit them under these budgets.

TSN2 is the second exception, and it cuts the other way. Over all 61 instances its mixed pool is the best arm (3.41 against the control’s 3.60); over the 35 instances where both plateau, the margin vanishes entirely (2.34 against 2.32). The cells TSN2 fails are its largest — all ten 50×20 instances and eight of the ten 30×20 — which are precisely where Section 7.3 locates its gain. The benefit seeding brings to the strongest solver is thus confined to cells where its budget did not suffice to converge, a materially weaker claim than the unrestricted table supports, and we report it as such.

Both conclusions survive the rule’s own parameters. Varying the threshold over $\{0.05, 0.1, 0.25\}\%$ and the required fraction of runs over $\{0.5, 0.8, 0.9, 1.0\}$, and re-deriving the intersection at each setting, the GA reversal holds in every combination: the mixed pool beats the control by between 0.13 and 0.72 points, and its half-strength variant by between 0.33 and 0.72, with intersections ranging from 8 to 44 instances. On TSN2 the difference stays within ± 0.14 points of zero throughout, so the disappearance of its advantage is not an artefact of where the threshold is placed either. The sensitivity is computed on the intersection rather than on the control-plateau subset, since that is the comparison the two conclusions rest on.

Unless stated otherwise, the tables that follow use all 61 instances, so that every solver is evaluated on the same instance set; results restricted to converged cells are reported wherever they change a conclusion, which is the GA and TSN2.

7.3. Final quality at the budget

Table 4 reports mean RPD over the 61 instances. Three regimes are immediately visible, and they do *not* order themselves by solver strength.

Table 4: Mean RPD (%) with respect to published lower bounds, over 61 instances and 30 runs per cell. Lower is better; the best arm per solver is in bold. A0 is the unseeded control.

Solver	A0	V2H	V2	MOR	GT	GP	MIX	MIXH
GA	7.52	7.66	8.05	9.70	8.47	7.99	7.58	7.25
ABCE3	10.51	9.51	9.39	10.71	10.20	9.68	9.37	9.43
fEABCLS	3.67	3.66	3.66	5.16	4.47	4.04	3.67	3.58
TSN2	3.60	3.64	3.62	4.25	3.99	3.67	3.41	3.36

ABCE3: large and uniform gains.. The elitist bee colony is the weakest solver in the portfolio (10.51% RPD unseeded) and benefits massively: MIX reaches 9.37%, V2 9.39% and V2H 9.51%, and even GP, a mid-ranked generator elsewhere, reaches 9.68%. This is the textbook case: the solver’s own diversification never reaches the quality already present in the seeds.

TSN2: a gain confined to the largest instances.. The tabu-search hybrid is the strongest solver, and the natural expectation — which our own pilot supported — is saturation. It does not saturate, but the effect must be stated carefully. MIX lowers RPD from 3.60% to 3.41%, removing 5.3% of the excess over the lower bound, and it is the only arm that improves this solver: every other pool is neutral or harmful. The gain is small in aggregate and it is *not* uniform, being almost entirely attributable to one size class, which Section 7.5 quantifies. The defensible claim is therefore not that seeding improves the strongest solver in general, but that it does so on large instances under a search budget that is tight relative to instance size.

fEABCLS: neutrality.. The memetic ABC lands within a hundredth of a point of its control under the three best generators (3.67% against MIX 3.67%, V2

3.66%, V2H 3.66%): its hill-climbing component reaches seed-level quality unaided within the budget. Neutrality here is a property of the *good* pools only — MOR costs it 1.49 points.

GA: harmed by every pool but one.. Six of the seven seeded arms leave the genetic algorithm worse than its 7.52% control, V2 clearly so at 8.05% and the mixed pool at parity (7.58%). The exception is the mixed pool at half strength, MIXH, at 7.25%: the only arm anywhere in the study that improves this solver. Section 7.8 takes that up, since it is the one result that separates the two mechanisms this paper proposes. The rest of the reading is conditional on the budget: the GA fails the plateau rule on 49 of 61 instances, more than any other solver, and where its control does converge the ordering reverses (Section 7.2).

Across all four solvers the generator ranking is stable in its extremes: MOR is the worst seeded arm and GT the second worst on every solver, while MIX holds the best Friedman rank among seeded arms on all four (Section 7.10). On mean RPD, MIX is the best seeded arm on three solvers and is within 0.01 points of the best on fEABCLS. The positions in between are not stable: GP outranks V2 on the GA but not on the other three. The learned policy is thus competitive with the best tuned heuristic without dominating it, and it is the mixed pool rather than any single generator that is uniformly at the top — which already suggests that *pool diversity* rather than raw seed quality drives the benefit. Section 7.9 tests that directly.

7.4. Effect sizes with the instance as unit

The experimental unit is the instance. Each arm is summarised per instance by the mean of its 30 runs; the quantity analysed is the per-instance difference from the control, giving $n = 61$ observations; one signed-rank test per (solver, arm) is applied to those differences, with Holm correction over the six arms of each solver; and a 95% confidence interval for the mean difference is obtained by bootstrapping instances (10^4 resamples). Table 5 reports the result. Alongside it, W/NS/L counts the instances on which the arm is significantly better, not

significantly different, or significantly worse under a per-instance signed-rank test, and $\hat{A}_{12} < 0.5$ favours seeding; these give the same comparison in a more directly readable form.

Three of the four solvers need a word beyond the table.

ABCE3 gains unambiguously.. MIX improves it by 1.14 percentage points of RPD, interval $[-1.36, -0.93]$, $p < 10^{-3}$; V2, V2H and GP likewise. Correcting over all 24 pre-specified contrasts of the study rather than within each solver changes nothing: the same 15 remain significant.

The GA is harmed by homogeneous pools, not by seeding.. The mixed pool's difference is +0.055 points, with an interval spanning zero and $p = 0.55$, while every *single-generator* pool is significantly harmful: MOR (+2.18), GT (+0.95), V2 (+0.53) and GP (+0.47). Read with Section 7.2, the picture is a solver damaged by homogeneous seeds under an insufficient budget.

The improvement on TSN2 does not reach significance.. This is the reading that changes most. The signed-rank test over the 61 per-instance differences gives $p = 0.065$ before correction and $p = 0.26$ after, so MIX cannot be declared better than the control by a test that weights every instance equally. The bootstrap interval for the mean difference, $[-0.35, -0.045]$, nonetheless excludes zero. The two statistics are not in conflict: the signed-rank test asks about the median difference, which is essentially zero because 51 of the 61 instances are unaffected, while the mean is pulled negative by the ten 50×20 instances that carry 94.8% of the effect (Table 6). We therefore state the TSN2 result as what it is — a small mean improvement concentrated in one size class, with an interval that excludes zero but without significance under an instance-weighted test — and not as a general improvement of the strongest solver. Establishing the latter would require either more instances in the affected regime or a budget under which more of them behave like 50×20 .

Table 5: Each seeded arm against the unseeded control, with the instance as the experimental unit ($n = 61$). Differences are in percentage points of RPD; negative favours seeding, and boldface marks the best of the six pre-specified arms. † MIXH is the exploratory arm of Section 7.8: it is excluded from the Holm family so that an arm chosen after the fact cannot alter the pre-specified results, and its p is unadjusted. W/NS/L and $\hat{A}_{12}$ summarise the run-level distributions within each instance and are descriptive; the difference, interval and p are the instance-level inference.

Solver	Arm	W/NS/L	$\hat{A}_{12}$	Mean diff.	95% CI	p (Holm)
GA	V2H	11/31/19	0.55	+0.14	$[-0.07, +0.36]$	0.358
	V2	9/25/27	0.61	+0.53	$[+0.24, +0.84]$	0.012*
	MOR	1/3/57	0.92	+2.18	$[+1.89, +2.48]$	<0.001***
	GT	4/16/41	0.74	+0.95	$[+0.67, +1.23]$	<0.001***
	GP	9/21/31	0.65	+0.47	$[+0.20, +0.73]$	0.002**
	MIX	13/29/19	0.53	+0.06	$[-0.18, +0.29]$	0.546
	MIXH †	20/30/11	0.42	-0.28	$[-0.47, -0.10]$	0.007**
ABCE3	V2H	32/29/0	0.30	-1.00	$[-1.22, -0.80]$	<0.001***
	V2	35/26/0	0.27	-1.11	$[-1.34, -0.89]$	<0.001***
	MOR	1/54/6	0.54	+0.21	$[+0.09, +0.33]$	0.005**
	GT	14/46/1	0.44	-0.30	$[-0.50, -0.13]$	0.005**
	GP	30/31/0	0.32	-0.83	$[-1.08, -0.59]$	<0.001***
	MIX	35/26/0	0.27	-1.14	$[-1.36, -0.92]$	<0.001***
	MIXH †	42/19/0	0.28	-1.07	$[-1.30, -0.87]$	<0.001***
fEABCLS	V2H	7/43/11	0.50	-0.02	$[-0.12, +0.09]$	1.000
	V2	7/44/10	0.49	-0.01	$[-0.12, +0.10]$	1.000
	MOR	0/5/56	0.89	+1.48	$[+1.24, +1.74]$	<0.001***
	GT	3/12/46	0.77	+0.80	$[+0.60, +1.00]$	<0.001***
	GP	7/27/27	0.64	+0.37	$[+0.20, +0.54]$	<0.001***
	MIX	12/38/11	0.50	-0.01	$[-0.13, +0.11]$	1.000
	MIXH †	16/34/11	0.45	-0.09	$[-0.22, +0.03]$	0.226
TSN2	V2H	5/48/8	0.52	+0.05	$[-0.05, +0.14]$	0.486
	V2	6/47/8	0.51	+0.02	$[-0.08, +0.13]$	0.486
	MOR	9/13/39	0.69	+0.65	$[+0.39, +0.91]$	<0.001***
	GT	9/20/32	0.64	+0.40	$[+0.18, +0.61]$	<0.001***
	GP	10/32/19	0.54	+0.07	$[-0.14, +0.25]$	0.456
	MIX	15/41/5	0.44	-0.19	$[-0.34, -0.04]$	0.261
	MIXH †	18/36/7	0.43	-0.24	$[-0.41, -0.08]$	0.033*

7.5. The effect scales with instance size

Aggregate figures conceal a strong interaction with problem size. Figure 1 plots the relative change in $\mathbb{E}[C_{\max}]$ under MIX seeding against instance class, ordered by number of operations (100 to 1000).




Figure 1: Relative change in $\mathbb{E}[C_{\max}]$ under MIX seeding with respect to the unseeded control, by instance class (ordered by number of operations). Negative values favour seeding. The benefit for ABCE3 and TSN2 grows with instance size, whereas the penalty for the GA persists.

Two trends emerge. For ABCE3 the gain deepens with size, from -0.21% on the smallest instance to -1.24% on 20×20 and -1.50% on 50×20 . For TSN2 the effect is flat and near zero up to 50×15 and then drops sharply to -1.06% on 50×20 — and on that class *every* seeded arm beats the control, including GP ($\Delta = -30.9$), MOR ($\Delta = -25.6$) and GT ($\Delta = -19.6$), all three of which are harmful everywhere else. The GA moves in the opposite direction, remaining between $+0.05\%$ and $+0.48\%$ on the four largest classes.

Table 6 makes the concentration explicit for TSN2, and it is the single most important caveat in this paper. The ten 50×20 instances account for 94.8% of the aggregate improvement; the remaining 51 instances contribute almost noth-

ing. Reporting only the pooled figure of Table 4 would therefore misrepresent the result as a general improvement of the strongest solver, which it is not.

Table 6: Decomposition of the MIX seeding effect on TSN2 by instance class. Δ is the mean signed change in $\mathbb{E}[C_{\max}]$ (negative favours seeding); the last column is each class’s share of the total improvement.

Class	n	$\mathbb{E}[C_{\max}]$ A0	$\mathbb{E}[C_{\max}]$ MIX	Δ (%)	Share
10×10	1	937.1	944.5	+0.78	—
15×15	10	1253.6	1254.3	+0.06	—
20×20	10	1668.8	1669.1	+0.01	—
30×15	10	1844.8	1843.6	−0.06	3.6%
30×20	10	2052.6	2051.4	−0.06	3.4%
50×15	10	2780.3	2779.2	−0.04	3.2%
50×20	10	2970.1	2938.4	−1.06	94.8%

The interpretation is a budget-starvation argument. On 50×20 instances (1000 operations) even a strong tabu search cannot, within the 900 s cap, recover from a random start the structure that a constructive heuristic supplies for free. Seeding does not make the search better; it removes a phase of the search that the budget cannot afford. This predicts that the benefit of seeding should be reported as a function of the budget-to-size ratio rather than as a single aggregate figure, and that published seeding gains obtained on small instances will not transfer.

7.6. Anytime behaviour

Final quality alone cannot distinguish a solver that is genuinely improved by seeding from one that is merely given a head start. Because per-solver budgets range from 60 s to 900 s across size classes, absolute time is not comparable across instances; we therefore normalise each trace to the fraction of its own budget consumed, the standard convention for aggregated anytime curves, and average over instances on a common grid so that no instance enters or leaves mid-curve.

Figure 2 shows the resulting profiles over all 61 instances, and Table 7 gives the values at five budget fractions.




Figure 2: Anytime behaviour over all 61 instances. Mean RPD with respect to the published lower bound against the fraction of the per-instance budget consumed, for all eight arms.

One caution applies to these curves. They record the best individual *in the population* at each sampling point rather than the solution the solver finally returns, and the last sample precedes termination by up to one sampling interval. In practice the two agree closely: the curve endpoint differs from the final quality of Table 4 by a median of 0.000% on ABCE3, 0.003% on fEABCLS and 0.029% on the GA. On TSN2 the curve ends 0.041% *below* the returned quality, because TS-N2 ranks internally by LEX2 while the curve tracks $\mathbb{E}[C_{\max}]$, so a schedule that is better on the midpoint need not be the one it returns. We nonetheless confine every cross-arm comparison to a single solver, where all arms are measured identically.

Within that scope, the aggregate view exposes a result the final-quality

Table 7: Mean RPD (%) over all 61 instances at five fractions of the per-instance budget. Values are resampled from execution traces and agree with Table 4 to within the trace sampling resolution.

Solver	Arm	10%	25%	50%	75%	100%
GA	A0	13.58	11.00	8.96	8.06	7.56
	V2H	12.91	11.78	9.70	8.50	7.73
	V2	12.70	11.92	10.29	9.02	8.14
	MOR	14.07	12.24	10.85	10.17	9.75
	GT	13.15	11.34	9.76	8.99	8.51
	GP	12.43	11.09	9.48	8.60	8.04
	MIX	12.10	11.08	9.32	8.32	7.64
	MIXH	12.21	10.74	8.78	7.82	7.29
ABCE3	A0	12.28	11.45	10.83	10.67	10.52
	V2H	10.83	10.10	9.68	9.61	9.52
	V2	10.78	9.97	9.55	9.48	9.40
	MOR	12.84	11.95	11.11	10.92	10.73
	GT	11.84	11.12	10.52	10.35	10.22
	GP	10.65	10.17	9.84	9.76	9.68
	MIX	10.23	9.76	9.49	9.46	9.38
	MIXH	10.30	9.84	9.60	9.53	9.44
fEABCLS	A0	14.84	5.60	4.32	3.89	3.68
	V2H	9.50	5.61	4.35	3.90	3.67
	V2	9.36	5.63	4.35	3.90	3.67
	MOR	13.15	6.59	5.72	5.39	5.17
	GT	11.77	5.96	5.03	4.70	4.48
	GP	9.21	5.46	4.60	4.26	4.05
	MIX	8.91	5.27	4.23	3.87	3.67
	MIXH	8.85	5.08	4.10	3.76	3.58
TSN2	A0	12.48	5.98	4.36	3.79	3.53
	V2H	8.70	5.63	4.29	3.81	3.59
	V2	8.41	5.53	4.25	3.80	3.57
	MOR	10.04	5.52	4.66	4.36	4.19
	GT	9.78	5.36	4.40	4.08	3.93
	GP	7.94	5.08	4.11	3.76	3.59
	MIX	7.89	5.05	3.94	3.55	3.36
	MIXH	8.02	4.90	3.80	3.48	3.30

tables cannot show: *at short budgets seeding is a large win for every solver in the portfolio, but by margins that differ by an order of magnitude between them.* At 10% of budget MIX cuts mean RPD from 14.84% to 8.91% on fEABCLS and from 12.48% to 7.89% on TSN2, but only from 12.28% to 10.23% on ABCE3 and from 13.58% to 12.10% on the GA — reductions of 40%, 37%, 17% and 11% in the excess over the lower bound, equivalent to 5.2%, 4.1%, 1.8% and 1.3% of makespan. The two largest short-budget gains therefore accrue to the two *strongest* solvers, which is the reverse of the ordering at the budget limit. Both are the solvers whose control spends its opening on an unseeded random population before local search can act: fEABCLS falls from 14.84% to 5.60% between a tenth and a quarter of its budget, and seeding buys precisely that opening. The practical reading is that the question “does seeding help?” is ill-posed unless the budget is specified, and that which solver benefits most depends on where in its own trajectory the budget cuts.

The GA panel isolates the reversal. Its seeded arms lead until roughly 20% of the budget, after which the control crosses and stays below to the end. This is the anytime signature of the consolidation failure analysed in Section 8.2: the seeded population is ahead while quality is inherited from the seeds, and falls behind as soon as progress must come from recombination. ABCE3 shows the opposite profile, a gap that opens immediately and never closes.

The generator ranking is also stable along the whole trajectory, not merely at the end. MOR is the uppermost seeded curve in all four panels and GT the next, so their disadvantage is not an artefact of where measurement stops: from a quarter of the budget onwards both are worse than the control on the two ABC variants. At a tenth of the budget they are still ahead of it, which says more about how poor an unseeded random population is at that point than about any merit of theirs. MIX is the lowest curve in three of four panels for essentially the entire run. Two observations qualify the ranking. First, GP is competitive with MIX at very short budgets — on TSN2 at 10% of budget GP reaches 7.94% against MIX’s 7.89% — so the advantage of the mixed pool *widens as the search proceeds* rather than being inherited from higher initial

seed quality. That is exactly what one expects if the operative variable is pool diversity rather than seed quality, and Section 7.9 confirms it directly. Second, the seeded *fraction* interacts with the solver, and not in one direction. On the GA, seeding half the population beats seeding all of it: the paired per-instance difference between V2H and V2 is -0.39 points of RPD, interval $[-0.54, -0.26]$, $p < 10^{-3}$. On ABCE3 the comparison reverses with the same confidence, $+0.12$ points, $[+0.05, +0.18]$, $p < 10^{-3}$, in favour of seeding the whole population; on fEABCLS and TSN2 the two are indistinguishable ($p = 0.58$ and $p = 0.35$). Retaining random material is therefore not a general recommendation. It helps the solver that cannot consolidate a fully seeded population within its budget, which is the same solver, and on our reading the same mechanism, as the reversal of Section 7.2. Section 7.8 shows that this lever is independent of pool composition and adds to it.

Figure 3 restricts the same analysis to the 50×20 class in absolute time, where the seeding effect on the strong solvers is concentrated.

Two features are specific to this class. ABCE3’s control flattens at approximately 11.2% while both seeded arms flatten near 9.7%, a gap present from the first seconds that never closes — a genuine improvement in converged quality rather than an acceleration. Its panel is visibly coarser than the others because its budget on this class is the 60s floor of Equation (1) while the trace is sampled every 5s: thirteen measurements per run, against 165 for fEABCLS. The steps are where the data are, not an artefact of the estimator. TSN2’s control exhibits a sharp step near 50s, when the tabu phase first engages, and thereafter descends steadily; MIX remains below it over the entire trajectory and is still below at the budget limit (approximately 3.3% against 4.1%). The separation is therefore not an artefact of where the budget stops.

One methodological consequence deserves emphasis. The GA result would have been reported as a *success* under any budget shorter than its crossover point, which is precisely the artefact documented in Section 5. Since seeding studies commonly fix a budget without verifying convergence, this is a plausible source of optimistic seeding claims in the literature, and it argues for reporting




Figure 3: Convergence on the 50×20 class (10 instances, 30 runs). Mean RPD with respect to the published lower bound against CPU time. Time axes differ because budgets are calibrated per solver and size class. Markers show the instants at which the traces are actually sampled, drawn where that resolution is coarse enough to matter.

anytime profiles rather than single-budget summaries.

7.7. Time to target

Anytime curves show *that* seeded arms lead early; time-to-target quantifies *by how much*. We report the two metrics of Section 5: (A) the time to come within 1% of the arm’s own final value, and (B) the time to reach the unseeded control’s final quality within +0.5%. Both are expressed as a fraction of the per-instance budget so that classes with different budgets can be pooled, and both are medians over the 61 instances. Metric B is reported together with the proportion of instances on which the arm reaches the target at all, since a median computed only over successes is meaningless without it.

Metric A behaves exactly as the pilot warned, and one figure shows why.

Table 8: Time to reach the unseeded control’s final quality (+0.5%), paired by instance, as a fraction of the per-instance budget. Three representative arms per solver; the full table is released with the data.

Solver	Arm	Reach	Earlier/later	p (sign)	RMTT	Δ RMTT [95% CI]
GA	A0	100%	—	—	0.62	—
	V2	56%	6/49	< 0.001	0.77	+0.16 [+0.12, +0.20]
	MOR	5%	0/59	< 0.001	0.99	+0.37 [+0.31, +0.44]
	MIX	77%	18/37	0.057	0.70	+0.09 [+0.04, +0.13]
	MIXH [†]	85%	35/19	0.009	0.62	0.00 [−0.05, +0.06]
ABCE3	A0	100%	—	—	0.38	—
	V2	100%	54/2	< 0.001	0.16	−0.22 [−0.28, −0.16]
	MOR	70%	13/40	0.004	0.60	+0.22 [+0.14, +0.30]
	MIX	100%	57/2	< 0.001	0.12	−0.26 [−0.31, −0.21]
	MIXH [†]	100%	60/0	< 0.001	0.11	−0.27 [−0.32, −0.22]
fEABCLS	A0	100%	—	—	0.44	—
	V2	89%	24/25	1.00	0.50	+0.06 [+0.01, +0.11]
	MOR	11%	0/59	< 0.001	0.94	+0.49 [+0.43, +0.56]
	MIX	89%	28/26	1.00	0.50	+0.06 [+0.01, +0.12]
	MIXH [†]	90%	37/17	0.002	0.46	+0.02 [−0.04, +0.09]
TSN2	A0	100%	—	—	0.52	—
	V2	90%	35/24	0.62	0.57	+0.05 [−0.00, +0.10]
	MOR	43%	19/41	0.047	0.74	+0.22 [+0.12, +0.31]
	MIX	93%	42/15	0.001	0.46	−0.06 [−0.11, +0.01]
	MIXH [†]	95%	46/13	< 0.001	0.43	−0.09 [−0.15, −0.02]

On the GA the earliest plateau belongs to MOR (0.49 of the budget against the control’s 0.60), and MOR is the worst generator on all four solvers: an arm that settles quickly on a poor value scores well *because* it is poor. We report the metric to document that it must not be used alone, and do not tabulate it further.

Metric B needs care of a different kind. An arm that never reaches the target on an instance has no time-to-target there, and dropping such instances would summarise only the successes — precisely discarding the cases where the arm does worst. The reach rate and the times must therefore enter the analysis together.

Two properties of the design constrain how. First, the arms are paired: each seeded arm and its control share the instance, the budget and the target, so a comparison that treats them as independent groups is not available. Second, and more restrictive, the target on an instance is the control’s own endpoint, so the control attains it on all 61 instances by construction. Its sample cannot be censored, whereas the seeded samples can. A log-rank test against such a reference contrasts a censored sample with one that is uncensorable by definition, and we therefore do not report one.

What survives both constraints is a paired analysis of the recorded times, with a non-reacher recorded at the budget, $T = 1.0$. The quantity the mean then estimates is the restricted mean time $E[\min(T, 1)]$, for which that recording is a definition rather than an approximation. The primary contrast is the sign test on the paired earlier/later counts, ties excluded, taken against the null proportion the design itself produces rather than against 0.5: because the target is drawn from the control’s own endpoint, a split-half calibration under a true null (Section 7.12) puts that proportion at 0.44–0.49 depending on the solver. The paired difference in restricted mean time (ΔRMTT) is reported alongside it with a bootstrap interval over instances. Both are Holm-corrected over the six pre-specified seeded arms of each solver. [†]MIXH is the exploratory arm of Section 7.8: it is excluded from the correction family and its p is unadjusted. No test is reported for the control, which attains its own target on all 61 instances

by construction; “reach” is the percentage of instances on which an arm attains it at all.

Under this treatment ABCE3 stands out. Its mixed pool reaches the control’s *final* quality on every instance and earlier on 57 of the 59 that are not ties, cutting restricted mean time from 0.38 to 0.12 of the budget; V2 and V2H do nearly as well. These are unambiguous “same quality, sooner” results, and they agree with the converged-quality finding for that solver. TSN2 joins it more modestly: its mixed pool is earlier on 42 instances and later on 15 ($p = 0.003$), and reaches the target on 93% of them against the control’s 100% by construction.

The other two solvers give the opposite answer. *Five of the GA’s six seeded arms are significantly slower* to the control’s own final quality; only the mixed pool escapes, and narrowly (18 earlier, 37 later, $p = 0.057$), though its restricted mean time is clearly worse ($+0.09$, $[+0.04, +0.13]$). That is the same pattern the final-quality analysis found on this solver: every single-generator pool harms it and the mixed pool does not. On fEABCLS the three competitive pools — V2H, V2 and MIX — are statistically indistinguishable from the control ($p = 1.00$ for V2 and MIX), while the three weak ones are significantly slower. Seeding buys that solver nothing in time.

Mean and sign agree almost everywhere. They part on three arms: MIX on the GA, V2H on fEABCLS and GT on TSN2, where ΔRMTT reaches significance and the sign test does not. In each the difference is concentrated in a minority of instances rather than being a general shift, which is exactly what the sign test is there to detect.

Ten arm–solver pairs are significantly *slower* to the target than their control, and MOR is the extreme of every one: it is slower on all four solvers and reaches the target on no instance at all on the GA and fEABCLS, where 59 of 61 comparisons go against it. GT is slower on two solvers and, revealingly, *faster* on ABCE3 — the solver that tolerates every pool. Both generators sit at the bottom of the quality ladder of Section 7.9, so this is the one place in the study where seed quality and outcome move together.

The contrast with a success-only tally is instructive. Computing medians over successes alone would have credited MOR on the GA with reaching the target at 71% of budget, a figure derived from the 15% of instances where it succeeds at all.

7.8. *Composition and fraction as separate factors: an exploratory follow-up*

This subsection is exploratory and we mark it as such. The seven arms above were fixed before the experiment ran. The eighth, MIXH, was not: it was specified after the first seven had been analysed, precisely because their outcome raised a question they could not answer, and it was run under the same protocol three days later. Holm correction over the six pre-specified seeded arms handles simultaneous testing within a table; it does not undo the fact that this arm and its hypothesis were chosen with the earlier results in view. What follows is therefore a follow-up study whose findings invite confirmation, not a pre-specified test.

The question is this. Two of the study’s claims about how seeding works were measured on different pools: *composition*, from the mixed pool against the single-generator ones at full strength, and *fraction*, from the half-seeded learned arm against the fully seeded one. Whether they are one effect seen twice or two that combine is what the arm at their intersection can address.

The four arms form a complete 2×2 : pool (v2 or mixed) crossed with fraction (all or half the population). On the genetic algorithm neither factor alone helps — half-seeding the learned pool gives +0.14 points of RPD ($p = 0.36$) and mixing at full strength +0.06 ($p = 0.55$) — while their combination does: MIXH reaches −0.28 points, interval $[-0.47, -0.10]$, $p = 0.022$, and beats each component directly, −0.33 against MIX and −0.42 against V2H, both $p < 10^{-3}$ paired. It is the only arm in the study that improves this solver at its budget.

A combination outperforming each of its parts is compatible with several models, so we estimate the interaction rather than assert one. The difference in differences, $(\text{MIXH} - \text{MIX}) - (\text{V2H} - \text{V2})$, is +0.06 points on the GA with interval $[-0.06, +0.17]$ and $p = 0.48$: the data are consistent with the two factors

combining additively and exclude any large interaction, but do not demonstrate additivity. On two other solvers the interaction is small and *negative* — -0.08 on fEABCLS ($[-0.15, -0.03]$, $p = 0.008$) and -0.07 on TSN2 ($[-0.14, -0.01]$, $p = 0.046$) — so there the combination is worth slightly more than the sum of its parts. On ABCE3 it is indistinguishable from zero. The defensible summary is that the two are separable design levers whose effects do not cancel, not that they are additive.

The time-to-target evidence points the same way and is stronger, since it survives correction. Against its control the mixed pool at full strength is *later* on this solver (18 earlier, 37 later); at half strength the sign inverts to 35 earlier and 19 later, $p = 0.013$, with restricted mean time falling from 0.70 to 0.62 — the control’s own value. Half seeding does not merely reduce the delay that full seeding causes here; it removes it.

Elsewhere the two mechanisms do not add, and on one solver they trade off. MIXH is the best arm on fEABCLS (3.58% against the control’s 3.67%) and on TSN2 (3.36% against 3.60%), and on both it converts a time-to-target result that failed correction into one that passes: fEABCLS moves from 28/26 at $p = 1.00$ to 37/17 at $p = 0.010$, and TSN2 from 42/15 to 46/13. But on ABCE3, the solver that consolidates seeds most readily, the full mixture remains better (9.37% against 9.43%; $+0.07$ in favour of MIX, $p = 0.096$), and it is the one solver where retaining random material buys nothing.

That partial seeding can beat full seeding is not new in itself: it has been reported for multi-objective combinatorial problems, where populations containing only a few heuristic solutions outperformed both random and heavily seeded ones [27]. What the present design adds is that the prescription does not transfer uniformly. Under an identical protocol the same reduction in fraction helps three solvers and costs the fourth, and the pre-specified fraction sub-study of Section 6, run on the solvers where it did not help, found the benefit monotone in the fraction with no interior optimum. The fraction is a solver-dependent lever, not a constant of the problem.

That pattern is the same one the convergence analysis produces. Retaining

random material helps exactly where the solver cannot consolidate a fully seeded population inside its budget, and the ordering of solvers by how much it helps is the ordering by how badly they consolidate. We report this as a coherent association across three independent measurements — converged quality, time to target, and the fraction contrast — and not as an isolated effect. Three cautions bound it. The arm and its hypothesis were selected after the first seven arms had been analysed, so this is a follow-up finding awaiting independent confirmation. The genetic algorithm’s final-quality result does not survive correction over all 28 contrasts of the study, although its time-to-target result does. And only two fractions were run, so the shape of the response between them is unmeasured; the 2×2 estimates an interaction at those two levels, not a curve.

7.9. Initial-population quality: the mixed pool does not start better

The obvious explanation for MIX’s advantage would be that its seeds are simply better. They are not. Table 9 reports the quality of the initial population each arm actually receives, read from the step-zero record of the traces, so it is the solvers’ own evaluation rather than the value stored with the pool.

Table 9: Quality of the seeded initial population relative to the unseeded control, averaged over five instances spanning 20×20 to 50×20 (TSN2 traces). Negative is better. “Mean” is over all 250 initial individuals and is the quantity that characterises the pool.

Arm	Best individual	Mean of the 250
V2H	−11.34%	−7.35%
V2	−11.84%	−14.73%
MOR	−5.28%	−10.40%
GT	−4.75%	−8.85%
GP	− 12.09%	− 15.48%
MIX	−12.25%	−13.12%

On the mean of the initial population — the natural measure of pool quality — MIX is *worse* than both V2 and GP, by 1.6 and 2.4 percentage points

respectively, and it never holds the best individual by a meaningful margin. It nonetheless finishes ahead of both (Table 4). A pool that starts worse and ends better is direct evidence that the operative variable is composition, not seed quality.

This also bears on a possible confound. The mixed pool introduces no new material: it is drawn from the same generator pools that feed the V2, GT and GP arms, taking the first 342/341/341 entries of each. For GT and GP, whose pools are i.i.d., that is a sample of the same distribution the single-generator arm draws from. For v2 it is not: those first 342 entries are one training checkpoint, while the V2 arm cycles through all three. So no difference in *provenance* explains the mixed pool’s advantage, but its v2 component is narrower than the V2 arm’s, not merely a re-mixture of it.

7.10. Which seeding strategy? Generator ranking and its solver-dependence

The practitioner’s question is not whether seeding helps in general but which pool to use, and whether that choice must be re-made for every solver. We answer both with a Friedman test over the eight arms with the 61 instances as blocks, followed by Holm-corrected paired post-hoc comparisons against the top-ranked arm of each solver. Table 10 reports the mean ranks.

Table 10: Mean Friedman ranks of the eight arms (1 = best) over 61 instances. Friedman rejects rank equality for every solver ($p < 10^{-15}$ throughout). Boldface marks arms not significantly worse than that solver’s top-ranked arm (Holm-corrected paired Wilcoxon, $\alpha = 0.05$).

Solver	A0	V2H	V2	MOR	GT	GP	MIX	MIXH
GA	3.39	3.80	5.05	7.36	5.67	4.64	3.75	2.34
ABCE3	6.55	3.61	2.99	7.09	5.88	3.98	2.87	3.03
fEABCLS	3.55	3.45	3.57	7.47	6.41	5.15	3.48	2.93
TSN2	4.06	4.58	4.45	6.25	5.98	4.67	3.13	2.87

Two findings follow, and they answer different questions.

Which generator: the answer varies little across solvers.. The ranking of the seven seeded arms is consistent across the portfolio: the mean pairwise Spearman correlation between the four rank vectors is $\rho = 0.72$. A mixed pool is the best seeded arm on every solver, at half strength on three of them (ranks 2.34 to 2.93) and at full strength on ABCE3 (2.87); the two variants are never separated by more than 0.16 of a rank except on the GA. MOR is last for all four (6.25 to 7.47) and GT second-to-last for three. With eight arms and four solvers a rank correlation of 0.72 is a descriptive summary rather than a demonstration of invariance, and the intermediate positions do move: GP outranks V2 on the GA but not elsewhere. What is stable is the top and the bottom, which is what a practitioner needs: within this portfolio, use a mixed pool and avoid the low-diversity constructive pools, without re-tuning that choice per algorithm.

Whether to seed at all: the answer does depend on the solver.. The only arm whose rank moves substantially across solvers is the unseeded control, from 3.39 on the GA, where only the half-strength mixed pool beats it, to 6.55 on ABCE3, where it ranks seventh of eight and is beaten by every seeded arm except MOR. The weakest pair of solvers is GA versus ABCE3 at $\rho = 0.40$, against 0.86–0.90 among the other pairs, and the displacement of A0 is what separates them.

This dissociation is the cleanest statement of the paper’s result. The question “which seeds are good?” has a solver-independent answer, because it is a property of the pool. The question “should I seed?” does not, because it depends on whether the solver’s variation operators can exploit the pool — the mechanism analysed in Section 8.2. Reporting only the first, as seeding studies commonly do, produces a recommendation that silently fails on solvers like the GA.

Why the advantage is not an artefact of sampling.. A natural objection is that MIX might win because it exposes the search to more distinct seed material. Three properties of the design rule this out. The five pools are the same size ($L = 1024$), so every arm draws 30 blocks of 250 under the same arithmetic. The mixed pool is assembled from precisely the solutions that feed the V2, GT and

GP arms, so it contains no material those arms do not also see. And directly measured (Section 7.9), the population MIX actually receives is on average *worse* than the ones V2 and GP receive, yet it converges to better solutions. Sampling breadth and seed quality both point the wrong way for the objection; within-population generator diversity is what remains.

7.11. *The pilot’s five statements, at scale*

Section 5 closed with five statements fixed before the final experiment. Table 11 reports their fate, which is part of the point of specifying them. Two are worth a sentence beyond the verdict. The GA statement is the one the pilot got right for the wrong reason: the gain does reverse under a converged budget, as it found on three instances, but on the cells where that solver’s control genuinely converges the reversal reverses again. And statement (v) is the one the experiment contradicts outright; it is the observation that reoriented the mechanism analysis of Section 8.

7.12. *Threats addressed by design*

Before the checks themselves, one inferential property of the time-to-target design needs establishing rather than assuming. Its target is *endogenous*: it is drawn from the control’s own endpoint, so the control is measured against a quantity taken from its own realisation while the seeded arm is not. That asymmetry is not obviously harmless, and it is not enough to observe that censoring cannot reverse a sign. We therefore measured it. Splitting each instance’s thirty control runs at random into two halves of fifteen, averaging each half and treating one as control and the other as arm, reproduces the analysis under a null that is true by construction. The null proportion is estimated from 500 such replicates and the resulting procedure validated on 500 further, independent ones, so that what is checked is the test actually performed rather than a different one. Under the null the seeded arm counts as earlier on 46.3%, 43.8%, 47.3% and 48.7% of instances on the four solvers, a tilt against seeding rather than the 50% a naive sign test would assume. Every sign test in Table 8 is taken

Table 11: The five statements fixed after the pilot, and their fate at scale. Section references give the supporting evidence.

Statement	Verdict	Evidence
(i) Seeding improves converged quality for the elitist ABC	confirmed, by a wider margin	35 wins, 0 losses of 61; every pool but MOR helps (§7.3)
(ii) With strong local search, quality saturates but metric-A acceleration remains	saturation confirmed, acceleration refuted	fEABCLS unchanged at the budget and indistinguishable under both time-to-target metrics; only the short-budget gain survives (§7.7)
(iii) For the plain GA the apparent gain is an artefact	confirmed at the pilot’s budgets, incomplete	reverses again where the control converges (§7.2)
(iv) GP and v2 tie at the top of the ladder; MOR never helps	partly confirmed	MOR is last on all four solvers, but the mixed pool outranks GP and v2 everywhere (§7.10)
(v) Initial-population quality ranks as the pool ladder predicts	refuted for the mixed pool	MIX starts worse than V2 or GP and still converges better (§7.9)

against its own solver’s proportion, and on the validation replicates that test rejects at 4.0%, 4.8%, 6.0% and 5.6% against a nominal 5%. Testing against 0.5 instead would reject at up to 8.4%, so the correction is what makes the procedure calibrated rather than anticonservative.

Two properties of the plateau rule of Section 7.2 qualify what it can be asked to support. Its *counts* are sensitive to how it is set: converged cells move from 114 at a 0.05% threshold to 162 at 0.25%, and from 194 to 63 as the fraction of runs required to satisfy it goes from half to all. Only the conclusions drawn from those counts are reported as stable, and only where they are. And on TS-N2, 325 of the 1830 control runs record a *negative* improvement, down to

−1.95%: the framework’s statistic takes the best individual under the solver’s own ranking, which for TS-N2 is LEX2, so the recorded incumbent’s midpoint need not be monotone even though its LEX2 value is. Those runs satisfy a rule that asks for improvement below 0.1% for a reason unrelated to convergence, which inflates that solver’s converged count by an amount we cannot separate cleanly.

Three further checks guard the numbers above. First, expected makespans were recomputed by an *independent* verifier: a standalone program that shares no code with the solver. It re-implements the instance parser, the interval arithmetic and the decoding from scratch, depends only on the standard library, and includes no header of the framework, so a defect in the solver’s scheduling or interval classes cannot hide behind it. The sample is systematic rather than random — every combination of the four solvers and eight arms on one instance from each of the seven size classes, 980 stored solutions in total — so each solver, arm and class is represented.

The verifier reproduces the reported interval makespan exactly in 979 of the 980 cases. This is the expected outcome once one accounts for Lamarckian replacement: after local search the framework re-encodes the improved schedule into the chromosome, so the stored sequence already carries the machine orders of the schedule that was actually executed, and decoding it returns that same schedule. The single exception, one run on `tai50_20_09`, differs by 0.5 units of $\mathbb{E}[C_{\max}]$ on a value of 3119 — 0.016%, and in the direction that makes the reported figure the more pessimistic. We have not traced its cause and record it as an open point. Second, every arm of a solver runs on the same instance with the same seed indices and the same budget, so the design is blocked by instance and no arm is systematically advantaged; the qualification of Section 6 on what that pairing does and does not guarantee applies. Third, the stopping rule is a fixed CPU-time budget per size class, identical across arms, which eliminates the stagnation-triggered artefact documented in Section 5.

8. Discussion

8.1. Seeding quality is not the operative variable

The pre-specified hypothesis — that seeding helps precisely those solvers whose own search saturates below the seed level — is not supported. The ordering by unseeded quality and the ordering by seeding effect do not match: the weakest solver gains most, the strongest gains a little, a mid-ranked one is unaffected and another is harmed at the budget and helped once it suffices. Solver quality alone does not predict the sign.

Two variables do carry information, and the design supports them as associations rather than isolating them as causes. The first is the *budget relative to what the solver needs*. Section 7.2 shows the GA reversal to be a budget effect, and Section 7.5 shows the TSN2 gain confined to the class where the budget is tightest relative to instance size. The two findings point the same way: seeding matters most where the solver has least time to build structure itself, and its apparent sign can flip when that time is short enough.

The second is pool composition. MIX — equal thirds of v2, GT and GP, i.e. 84/83/83 of the 250 seeded individuals — holds the best Friedman rank among seeded arms on all four solvers, while the single-generator pools of comparable individual quality (V2, GP) are weaker and the low-diversity pools (MOR, GT) are actively harmful. Since MIX is assembled from exactly the solutions that feed V2, GT and GP, and starts from a worse average population than either V2 or GP (Section 7.9), neither provenance nor seed quality can account for its advantage. What we cannot do is separate diversity from the other properties a mixture has — coverage of distinct basins, complementarity between specific pairs of generators — because we did not run pools matched on quality with diversity varied independently. The composition claim is therefore an association supported by an elimination of the obvious alternatives, not a controlled result.

8.2. Why the genetic algorithm is harmed: a failure to consolidate

The standard explanation for a seeding-induced loss is premature convergence: a population of good, similar individuals collapses onto a local optimum.

It is the mechanism Burke et al. identified in evolutionary timetabling, where a seeding procedure returning similar solutions cost more in diversity than it gained in quality [24], and it is the first hypothesis to test here. The data do not support it. We instrumented the GA and ABCE3 with the mean pairwise Hamming distance of the population, sampled every 10 s over 10 runs on four medium instances (normalised to $[0,1]$, where 0 denotes an identical population). Table 12 reports the result. The seeded GA is not less diverse than its control — it is *more* diverse, by +20.1% under V2 and +10.7% under MIX averaged over the four instances, in the same direction on all four of them, and it ends the run at 0.720 against the control’s 0.464. We attach no p -value: with the instance as the unit and four instances, the smallest attainable two-sided sign-test p is 0.125, so this evidence is the size of the effect and its consistency, not a significance claim. ABCE3, which does not recombine, shows nothing in either direction (−1.4% under V2, +0.4% under MIX, two and one instances agreeing in sign).

The diversity-collapse account therefore fails on its own terms here: the arm that loses ground is the more diverse one, not the less. Whatever costs the seeded GA its advantage, it is not the mechanism of [24].

Table 12: Mean pairwise Hamming distance of the population (0 = identical population, 1 = maximal positional difference), averaged over four instances and 10 runs. The seeded GA converges structurally *less* than its control; ABCE3 is unaffected by seeding.

Solver	Arm	Start	Mean	End	Drop
GA	A0	0.9474	0.7646	0.4639	−51.0%
	V2	0.9310	0.8845	0.7195	−22.7%
	MIX	0.9301	0.8350	0.5552	−40.3%
ABCE3	A0	0.9474	0.7140	0.8541	−9.9%
	V2	0.9310	0.7059	0.8602	−7.6%
	MIX	0.9301	0.7171	0.8497	−8.6%

The failure is therefore the opposite of premature convergence: within a

fixed budget the seeded GA never consolidates. Its control contracts by 51.0% in structural spread and exploits the region it has found; under V2 the seeded variant contracts by only 22.7% and is still exploring when the budget expires. The proximate cause is that the seeds, being drawn from generators with different construction logic, are good but mutually *dissimilar*; JOX recombining two structurally divergent parents produces offspring that inherit the operation ordering of neither, so selection cannot accumulate structure. The harm grows with instance size (Section 7.5) because longer chromosomes need more generations to consolidate, and the budget does not grow proportionally.

The contrast with ABCE3 is consistent with this reading. Seeding does not perturb ABCE3’s diversity trajectory at all (differences of -2.3% and $+4.0\%$, $p = 0.27$ and $p = 0.90$), and its population stays dispersed throughout (0.85 at the end). ABCE3 does not rely on recombining population members to build structure, so seeds are retained and refined rather than destroyed.

The ordering of the two seeded arms is itself informative. V2, drawn from a single generator, produces the least consolidation (-22.7%) and the largest loss in final quality; MIX, the more heterogeneous pool, consolidates further (-40.3%) and is the better of the two arms on this solver. Structural heterogeneity of the *pool* and of the *surviving population* are therefore not the same quantity: what slows the GA is a population its operator cannot consolidate, not a diverse pool as such.

Three limits on this account should be stated plainly. First, it is consistent with the convergence analysis rather than independent of it: Section 7.2 shows the GA’s loss disappears where its control converges, which is what a consolidation-rate explanation predicts, but both observations rest on the same runs. Second, GA and ABCE3 differ in selection, replacement, mutation and population dynamics as well as in recombination, so attributing the contrast to JOX specifically is under-determined; establishing it would require holding the rest of the GA fixed and varying the crossover, which we did not do. Third, mean pairwise Hamming distance is a genotypic measure on a representation with repeated symbols, and need not track the building blocks JOX actually

manipulates; a measure of parent–offspring precedence preservation would test the mechanism more directly. We therefore offer consolidation failure as the reading the present data favour, not as a demonstrated cause.

8.3. Practitioner guidance

The results support four recommendations, each with its scope attached.

(i) Seed with a *mixed* pool. It outranks every single-generator pool on all four solvers, and drawing on three generators costs no more than drawing on one since the pool is the same size. Whether to replace the whole population or half of it is the open question: our half-strength arm is the best performer on three solvers and the only one that helps the genetic algorithm, but it is the exploratory arm of Section 7.8, so we offer it as the variant worth trying first rather than as an established default. What is not free is producing the pool: at the tenth of the budget where the advantage is largest, generating a run’s seeds costs several times that slice unless the pool is reused across runs (Section 4.4). Seed once, solve many.

(ii) Expect the benefit to grow as the budget tightens relative to instance size, and expect its *sign* to be unreliable when the budget is far from sufficient — that is where our genetic algorithm reverses.

(iii) Before concluding that seeding fails on a recombination-based solver, check whether its control has converged; on our data that distinction separates a genuine null from an artefact of the budget.

(iv) Do not assume a seeded fraction transfers between solvers. Seeding the whole population is better on the elitist ABC and worse on the genetic algorithm, both at $p < 10^{-3}$, and neither is detectable without running the contrast: the half-seeded and fully seeded arms differ by at most 0.39 points of RPD, well inside the spread that separates the generators themselves.

8.4. Limitations

Scope. The study covers one problem family (interval JSP), one encoding (permutation with repetition), and four solvers three of which share the ABC framework or components of it. Nothing here licenses a general statement about recombinaive metaheuristics or about seeding at large; the conclusions are about

this problem, this representation and these operators. Replication on the deterministic JSP, on a different uncertainty model, or on an external benchmark would be the natural next test.

Mechanism.. The two explanatory claims — pool composition rather than seed quality, and consolidation failure rather than premature convergence — are associations that survive the obvious alternatives, not controlled results. Section 8.2 states the three specific limits on the second of them.

Budgets and cost.. Budgets are CPU-time caps measured with the process clock on a single Linux machine at a fixed occupancy of fourteen concurrent workers. Within a cell every arm shares the cap, so the paired comparisons that carry every claim here are unaffected, but the absolute level is not portable: the work a CPU-second buys depends on contention for cache and memory bandwidth among the concurrent workers. Our absolute RPDs should not be compared against figures published for these solvers elsewhere, and more generally a single absolute quality figure reported under a time budget, without the hardware occupancy that produced it, is less reproducible than it appears. Section 4.4 reports the cost of producing the pools, with two gaps: the training of the learned policy is excluded as amortised across instances, and only the two heuristic generators were timed directly, the learned policy’s pool cost being inferred from per-rollout inference and MOR- ε ’s not measured at all.

The v2 pool is blocked, not i.i.d.. Its 1,024 entries are three consecutive blocks, one per training checkpoint, and the blocks differ in quality. Because an arm draws a contiguous window whose start advances with the run index, the checkpoint mixture varies from run to run: several V2 runs and twenty of the thirty V2H runs are fed by a single checkpoint. Over thirty runs the mixture averages close to balanced (91/80/80 of 250 for V2), so the effect on the per-instance means that carry our analysis is small, but the between-run variance of the two v2 arms is inflated relative to every other arm, and inflated in a structured rather than random way. The mixed pool inherits the head of that pool, so

its v2 third is one checkpoint. Interleaving the checkpoints, as the mixed pool interleaves its generators, would remove this.

Design gaps.. The seeded fraction $p^* = 100\%$ was selected on two solvers and eight instances, so it is a protocol decision for this study rather than a generally preferable fraction, and the half-seeded arm shows the fraction interacting with the solver in both directions (Section 7.6). Only two fractions were run, so the shape of that interaction between them is unmeasured. The mixed pool was not compared against two-generator mixtures or against alternative proportions, so we cannot attribute its advantage to any particular pairing. And the structural-diversity instrumentation covers four instances and two solvers, a narrow base for the argument it supports.

9. Conclusions

We conducted, to our knowledge, the largest controlled study of population seeding for the interval job shop: four solvers, eight initialisation arms, 61 benchmark instances and 30 runs per cell — 58,560 runs — evaluated against published lower bounds under a fixed-budget protocol designed to eliminate three artefacts we document in a pilot study.

Seeding is not uniformly beneficial, and its sign is not predicted by solver strength. Analysed with the instance as the unit, a mixed pool improves the weakest solver decisively (-1.14 points of RPD, 95% CI $[-1.36, -0.93]$), leaves a memetic ABC unchanged, and does not harm a genetic algorithm, though every single-generator pool does. On the strongest solver the mean improvement is small, concentrated in one size class, short of significance under an instance-weighted test, and absent altogether on the instances where that solver’s own control converges; we report it as a bounded result rather than as an improvement of the state of the art.

Three findings seem to us to travel beyond this problem family. First, *the budget decides the sign, not just the size*. The two effects at the ends of our

portfolio both move when the analysis is restricted to cells whose control actually converged, and they move in opposite directions: the genetic algorithm’s apparent damage disappears, and the tabu hybrid’s gain disappears with it. Its benefit is confined to the largest instances, which are exactly the cells where its budget did not suffice to converge. Neither reversal is visible in a table that pools converged and budget-limited cells. A seeding study that fixes one budget without checking convergence can therefore report either sign, and which one it reports is not predictable from the solver’s strength.

Second, *what a pool is made of, and how much of the population it replaces, are two separable levers*. The mixed pool draws on the solutions that feed the single-generator arms and receives a worse average initial population than two of them, yet outranks all of them: composition, not seed quality, orders the pools that help. Independently of that, the fraction seeded interacts with the solver, and in both directions — seeding half the population is better on the genetic algorithm and worse on the elitist ABC. The two levers do not cancel: neither alone improves the genetic algorithm and their combination does, and that combination is the best arm on three of the four solvers, although it was specified after the first seven and awaits confirmation. Seed quality is not irrelevant at the bottom of the ladder — MOR, the weakest generator, is significantly slower to the control’s own final quality on all four solvers — but among the pools that help it is not what separates them.

Third, *the failure mode is not the one usually assumed*. Premature convergence is the standard explanation for seeding going wrong; our diversity measurements show the opposite, a population that stays dispersed and never consolidates within the budget. We advance consolidation failure as the reading the data favour, while noting that separating it from the other differences between the two solvers would need ablations this design does not contain.

The practical consequence is that seeding should be treated as an interaction between pool composition, variation operator and budget, evaluated at the instance size of the intended application, rather than as a property of the seeds alone. And the methodological consequence is that the three habits our pilot

identified — stagnation-triggered stopping, single-budget comparison, and self-referential speed metrics — are enough to manufacture a positive seeding result where none exists, or to hide one that is there.

Data and code availability

The seed pools, the injection operator, the experiment setups, the run logs of all 58,560 runs, the aggregated result tables and the analysis scripts that produce every table and figure in this paper, together with the recovered crisp instance matrices of Section 7.1 and their hashes, are deposited at [DOI pending], together with the independent verifier of Section 7.12 and the commit hash of the solver framework.

References

- [1] D. Lei, Population-based neighborhood search for job shop scheduling with interval processing time, *Computers & Industrial Engineering* 61 (2011) 1200–1208. doi:10.1016/j.cie.2011.07.010.
- [2] D. Lei, Interval job shop scheduling problems, *International Journal of Advanced Manufacturing Technology* 60 (2012) 291–301. doi:10.1007/s00170-011-3600-3.
- [3] H. Díaz, I. González-Rodríguez, J. J. Palacios, I. Díaz, C. R. Vela, A genetic approach to the job shop scheduling problem with interval uncertainty, in: *Information Processing and Management of Uncertainty in Knowledge-Based Systems (IPMU 2020)*, Vol. 1238 of CCIS, Springer, 2020, pp. 663–676. doi:10.1007/978-3-030-50143-3_52.
- [4] H. Díaz, J. J. Palacios, I. González-Rodríguez, C. R. Vela, Elite artificial bee colony for makespan optimisation in job shop with interval uncertainty, in: *Bio-inspired Systems and Applications: from Robotics to Ambient Intelligence (IWINAC 2022)*, Vol. 13259 of LNCS, Springer, 2022, pp. 98–108. doi:10.1007/978-3-031-06527-9_10.

- [5] H. Díaz, J. J. Palacios, I. González-Rodríguez, C. R. Vela, An elitist seasonal artificial bee colony algorithm for the interval job shop, *Integrated Computer-Aided Engineering* (2023). doi:10.3233/ICA-230705.
- [6] H. Díaz, J. J. Palacios, I. González-Rodríguez, C. R. Vela, Fast elitist ABC for makespan optimisation in interval JSP, *Natural Computing* 22 (2023) 645–657. doi:10.1007/s11047-023-09953-2.
- [7] H. Díaz, Neighbourhood structures and ranking operators for the interval job shop scheduling problem, under review, *Computers & Operations Research* (2026).
- [8] C. Zhang, W. Song, Z. Cao, J. Zhang, P. S. Tan, X. Chi, Learning to dispatch for job shop scheduling via deep reinforcement learning, in: *Advances in Neural Information Processing Systems*, Vol. 33, 2020, pp. 1621–1632.
- [9] Y. Bengio, A. Lodi, A. Prouvost, Machine learning for combinatorial optimization: a methodological tour d’horizon, *European Journal of Operational Research* 290 (2) (2021) 405–421.
- [10] B. Kazimipour, X. Li, A. K. Qin, A review of population initialization techniques for evolutionary algorithms, in: *IEEE Congress on Evolutionary Computation (CEC)*, 2014, pp. 2585–2592.
- [11] S. J. Louis, J. McDonnell, Learning with case-injected genetic algorithms, *IEEE Transactions on Evolutionary Computation* 8 (4) (2004) 316–328.
- [12] P. J. M. van Laarhoven, E. H. L. Aarts, J. K. Lenstra, Job shop scheduling by simulated annealing, *Operations Research* 40 (1992) 113–125.
- [13] J. Carlier, E. Pinson, An algorithm for solving the job-shop problem, *Management Science* 35 (2) (1989) 164–176, proves optimality of the Fisher–Thompson 10x10 instance (930).
- [14] E. Nowicki, C. Smutnicki, A fast taboo search algorithm for the job shop problem, *Management Science* 42 (6) (1996) 797–813.

- [15] E. Nowicki, C. Smutnicki, An advanced tabu search algorithm for the job shop problem, *Journal of Scheduling* 8 (2) (2005) 145–159. doi:10.1007/s10951-005-6364-5.
- [16] E. Balas, A. Vazacopoulos, Guided local search with shifting bottleneck for job shop scheduling, *Management Science* 44 (2) (1998) 262–275. doi:10.1287/mnsc.44.2.262.
- [17] I. González-Rodríguez, C. R. Vela, J. Puente, R. Varela, A new local search for the job shop problem with uncertain durations, in: *Proceedings of ICAPS 2008*, AAAI Press, 2008, pp. 124–131.
- [18] J. J. Palacios, J. Puente, C. R. Vela, I. González-Rodríguez, Benchmarks for fuzzy job shop problems, *Information Sciences* 329 (2016) 736–752. doi:10.1016/j.ins.2015.09.042.
- [19] S. Afsar, C. R. Vela, J. J. Palacios, I. González-Rodríguez, Mathematical models and benchmarking for the fuzzy job shop scheduling problem, *Computers & Industrial Engineering* 183 (2023) 109454. doi:10.1016/j.cie.2023.109454.
- [20] J. O. Agushaka, A. E. Ezugwu, Initialisation approaches for population-based metaheuristic algorithms: a comprehensive review, *Applied Sciences* 12 (2) (2022) 896. doi:10.3390/app12020896.
- [21] M. Sarhani, S. Voß, R. Jovanovic, Initialization of metaheuristics: comprehensive review, critical analysis, and research directions, *International Transactions in Operational Research* 30 (6) (2023) 3361–3397. doi:10.1111/itor.13237.
- [22] H. Maaranen, K. Miettinen, M. M. Mäkelä, Quasi-random initial population for genetic algorithms, *Computers & Mathematics with Applications* 47 (12) (2004) 1885–1895. doi:10.1016/j.camwa.2003.07.011.

- [23] S. Rahnamayan, H. R. Tizhoosh, M. M. A. Salama, Opposition-based differential evolution, *IEEE Transactions on Evolutionary Computation* 12 (1) (2008) 64–79. doi:10.1109/TEVC.2007.894200.
- [24] E. K. Burke, J. P. Newall, R. F. Weare, Initialization strategies and diversity in evolutionary timetabling, *Evolutionary Computation* 6 (1) (1998) 81–103. doi:10.1162/evco.1998.6.1.81.
- [25] T. Friedrich, M. Wagner, Seeding the initial population of multi-objective evolutionary algorithms: a computational study, *Applied Soft Computing* 33 (2015) 223–230. doi:10.1016/j.asoc.2015.04.043.
- [26] C. Gong, Y. Nan, L. M. Pang, Q. Zhang, H. Ishibuchi, Effects of including optimal solutions into initial population on evolutionary multi-objective optimization, in: *Proceedings of the Genetic and Evolutionary Computation Conference (GECCO 2023)*, ACM, 2023, pp. 661–669. doi:10.1145/3583131.3590515.
- [27] C. Gong, Y. Nan, L. M. Pang, H. Ishibuchi, Q. Zhang, Initial populations with a few heuristic solutions significantly improve evolutionary multi-objective combinatorial optimization, in: *2023 IEEE Symposium Series on Computational Intelligence (SSCI)*, IEEE, 2023, pp. 1398–1405. doi:10.1109/SSCI52147.2023.10371937.
- [28] J. Park, J. Chun, S. H. Kim, Y. Kim, J. Park, Learning to schedule job-shop problems: representation and policy learning using graph neural network and reinforcement learning, *International Journal of Production Research* 59 (11) (2021) 3360–3377. doi:10.1080/00207543.2020.1870013.
- [29] W. Song, X. Chen, Q. Li, Z. Cao, Flexible job-shop scheduling via graph neural network and deep reinforcement learning, *IEEE Transactions on Industrial Informatics* 19 (2) (2023) 1600–1610. doi:10.1109/TII.2022.3189725.

- [30] J. Kotary, F. Fioretto, P. Van Hentenryck, B. Wilder, End-to-end constrained optimization learning: a survey, in: Proceedings of the 30th International Joint Conference on Artificial Intelligence (IJCAI), 2021, pp. 4475–4482. doi:10.24963/ijcai.2021/610.
- [31] F. Zhang, Y. Mei, S. Nguyen, M. Zhang, Survey on genetic programming and machine learning techniques for heuristic design in job shop scheduling, IEEE Transactions on Evolutionary Computation 28 (1) (2024) 147–167. doi:10.1109/TEVC.2023.3255246.
- [32] J. Ma, W. Gao, W. Tong, A deep reinforcement learning assisted adaptive genetic algorithm for flexible job shop scheduling, Engineering Applications of Artificial Intelligence 149 (2025) 110447. doi:10.1016/j.engappai.2025.110447.
- [33] M. López-Ibáñez, T. Stützle, Automatically improving the anytime behaviour of optimisation algorithms, European Journal of Operational Research 235 (3) (2014) 569–582. doi:10.1016/j.ejor.2013.10.043.
- [34] J. Derrac, S. García, D. Molina, F. Herrera, A practical tutorial on the use of nonparametric statistical tests as a methodology for comparing evolutionary and swarm intelligence algorithms, Swarm and Evolutionary Computation 1 (1) (2011) 3–18.
- [35] A. Vargha, H. D. Delaney, A critique and improvement of the CL common language effect size statistics of McGraw and Wong, Journal of Educational and Behavioral Statistics 25 (2) (2000) 101–132.
- [36] M. López-Ibáñez, J. Dubois-Lacoste, L. Pérez Cáceres, M. Birattari, T. Stützle, The irace package: iterated racing for automatic algorithm configuration, Operations Research Perspectives 3 (2016) 43–58.
- [37] H. Bustince, J. Fernandez, A. Kolesárová, R. Mesiar, Generation of linear orders for intervals by means of aggregation functions, Fuzzy Sets and Systems 220 (2013) 69–77.

- [38] É. Taillard, Benchmarks for basic scheduling problems, *European Journal of Operational Research* 64 (2) (1993) 278–285.
- [39] C. Bierwirth, A generalized permutation approach to job shop scheduling with genetic algorithms, *OR Spektrum* 17 (1995) 87–92.
- [40] M.-E. Coupvent des Graviers, L. Kobrosly, C. Guettier, T. Cazenave, Updating lower and upper bounds for the job-shop scheduling problem test instances, consolidated bounds for Taillard’s ta01–ta80 (Table 14) (2025). [arXiv:2504.16106](https://arxiv.org/abs/2504.16106).